

## ▼ Loading the dataset

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
d = pd.read_csv('/content/abalone.csv')
```

## ▼ Perform visualization

### Univariate analysis

```
d.head()
```


	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395

```
plt.hist(d['Diameter'])
```

```
(array([ 13.,  66., 180., 344., 513., 812., 1017., 934., 275.,
        23.]),
 array([0.055, 0.1145, 0.174, 0.2335, 0.293, 0.3525, 0.412, 0.4715,
        0.531, 0.5905, 0.65 ]),
 <a list of 10 Patch objects>)
```


```
plt.plot(d['Diameter'].head(10))
```

```
[<matplotlib.lines.Line2D at 0x7f079f97c790>]
```



```
plt.pie(d['Diameter'].head(), autopct='%0.2f')
```


```
([<matplotlib.patches.Wedge at 0x7f079f890150>,
 <matplotlib.patches.Wedge at 0x7f079f890910>,
 <matplotlib.patches.Wedge at 0x7f079f8981d0>,
 <matplotlib.patches.Wedge at 0x7f079f898ad0>,
 <matplotlib.patches.Wedge at 0x7f079f8a4650>],
 [Text(0.8507215626110557, 0.6973326486753676, ''),
 Text(-0.32611344931648134, 1.0505474849691026, ''),
 Text(-1.0998053664078908, -0.02069193128747144, ''),
 Text(-0.08269436219656089, -1.096887251480709, ''),
 Text(0.9758446362287218, -0.5076684409569241, ')],
 [Text(0.46402994324239394, 0.3803632629138369, '21.86'),
 Text(-0.17788006326353525, 0.5730259008922377, '15.87'),
 Text(-0.5998938362224858, -0.011286507974984419, '25.15'),
 Text(-0.045106015743578656, -0.5983021371712958, '21.86'),
 Text(0.5322788924883937, -0.2769100587037768, '15.27')])
```



## ▼ Bivariate Analysis


```
plt.scatter(d['Diameter'].head(500),d['Length'].head(500))
```

<matplotlib.collections.PathCollection at 0x7f079f869fd0>




```
sns.barplot(d['Sex'], d['Rings'])
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f079f7be650>




```
sns.barplot(d['Diameter'].head(50),d['Rings'].head(50))
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f079f7b0290>



```
sns.lineplot(d['Diameter'].head(),d['Rings'].head())
```


<matplotlib.axes.\_subplots.AxesSubplot at 0x7f079f7b3990>



## ▼ Multivariate Analysis


```
sns.boxplot(d['Sex'].head(10),d['Diameter'].head(10),d['Rings'].head(10))
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f079fef4550>



```
sns.pairplot(d.head(),hue='Rings')
```

&lt;seaborn.axisgrid.PairGrid at 0x7f079f382a50&gt;



## ▼ Perform the descriptive statistics on the dataset

```
d.head()
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7



```
d.tail()
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
4172	F	0.565	0.450	0.165	0.8870	0.3700	0.2390	0.2490	11
4173	M	0.590	0.440	0.135	0.9660	0.4390	0.2145	0.2605	10
4174	M	0.600	0.475	0.205	1.1760	0.5255	0.2875	0.3080	9
4175	F	0.625	0.485	0.150	1.0945	0.5310	0.2610	0.2960	10
4176	M	0.710	0.555	0.195	1.9485	0.9455	0.3765	0.4950	12



```
d.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4177 entries, 0 to 4176
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Sex                    4177 non-null   object
1   Length                 4177 non-null   float64
2   Diameter               4177 non-null   float64
3   Height                 4177 non-null   float64
4   Whole weight           4177 non-null   float64
5   Shucked weight         4177 non-null   float64
6   Viscera weight         4177 non-null   float64
7   Shell weight           4177 non-null   float64
8   Rings                  4177 non-null   int64
dtypes: float64(7), int64(1), object(1)
memory usage: 293.8+ KB
```

```
d.mode().T
```

	0	1
Sex	M	NaN
Length	0.55	0.625
Diameter	0.45	NaN
Height	0.15	NaN
Whole weight	0.2225	NaN
Shucked weight	0.175	NaN
Viscera weight	0.1715	NaN
Shell weight	0.275	NaN
Rings	9.0	NaN



```
d.shape
```

```
(4177, 9)
```

```
d.skew()
```

```
Length      -0.639873
Diameter    -0.609198
Height       3.128817
Whole weight 0.530959
Shucked weight 0.719098
Viscera weight 0.591852
Shell weight 0.620927
Rings        1.114102
dtype: float64
```

```
d.nunique()
```

```
Sex          3
Length       134
Diameter     111
Height       51
Whole weight 2429
Shucked weight 1515
Viscera weight 880
Shell weight 926
Rings        28
dtype: int64
```

## ▼ Check for Missing values and deal with them

```
d.isna()
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	False	False	False	False	False	False	False	False	False

```
d.isna().any()
```

```
Sex          False
Length       False
Diameter     False
Height       False
Whole weight False
Shucked weight False
Viscera weight False
Shell weight False
Rings        False
dtype: bool
```

```
True  False  False  False  False  False  False  False  False  False
```

```
d.isna().sum()
```

```
Sex          0
Length       0
Diameter     0
Height       0
Whole weight 0
Shucked weight 0
Viscera weight 0
Shell weight 0
Rings        0
dtype: int64
```

```
d.isna().any().sum()
```

```
0
```

## ▼ Find the outliers and Replace the Outlier

```
sns.boxplot(d['Diameter'])
```



<matplotlib.axes.\_subplots.AxesSubplot at 0x7f079b641210>

```
qnt=d.quantile(q=[0.25,0.75])
qnt
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
<b>0.25</b>	0.450	0.35	0.115	0.4415	0.186	0.0935	0.130	8.0
<b>0.75</b>	0.615	0.48	0.165	1.1530	0.502	0.2530	0.329	11.0

```
iqr=qnt.loc[0.75]-qnt.loc[0.25]
iqr
```

```
Length          0.1650
Diameter         0.1300
Height          0.0500
Whole weight     0.7115
Shucked weight   0.3160
Viscera weight   0.1595
Shell weight     0.1990
Rings           3.0000
dtype: float64
```

```
lower=qnt.loc[0.25]-(1.5*iqr)
lower
```


```
Length          0.20250
Diameter         0.15500
Height          0.04000
Whole weight    -0.62575
Shucked weight  -0.28800
Viscera weight  -0.14575
Shell weight    -0.16850
Rings           3.50000
dtype: float64
```

```
d['Diameter']=np.where(d['Diameter']<0.155,0.4078,d['Diameter'])
sns.boxplot(d['Diameter'])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f079b74cd10>
```

```
sns.boxplot(d['Length'])
```


```
<matplotlib.axes._subplots.AxesSubplot at 0x7f0799d3d490>
```



```
d['Length']=np.where(d['Length']<0.23,0.52, d['Length'])
```

```
sns.boxplot(d['Length'])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f0799cbb850>
```




```
sns.boxplot(d['Height'])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f0799c99ed0>
```




```
d['Height']=np.where(d['Height']<0.04,0.139, d['Height'])
d['Height']=np.where(d['Height']>0.23,0.139, d['Height'])
```




```
sns.boxplot(d['Height'])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f0799c0f350>
```



```
sns.boxplot(d['Whole weight'])
```


```
<matplotlib.axes._subplots.AxesSubplot at 0x7f0799be3250>
```



```
d['Whole weight']=np.where(d['Whole weight']>0.9,0.82, d['Whole weight'])
```

```
sns.boxplot(d['Shucked weight'])
```


<matplotlib.axes.\_subplots.AxesSubplot at 0x7f0799b69ed0>



```
d['Shell weight']=np.where(d['Shell weight']>0.61,0.2388, d['Shell weight'])
```

```
sns.boxplot(d['Shell weight'])
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f0799ac9610>



## ▼ Check for Categorical column and perform encoding

```
d['Sex'].replace({'M':1, 'F':0, 'I':2},inplace=True)
d
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	1	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.1500	15
1	1	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.0700	7
2	0	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.2100	9
3	1	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.1550	10

- Split the data into dependent and independent variables

[ ] ↪ 2 cells hidden

- Scale the independent variables

[ ] ↪ 2 cells hidden

4177 rows × 9 columns

- Split the Data into training and testing

[ ] ↪ 2 cells hidden

- Build the model

[ ] ↪ 1 cell hidden

- Train the model

[ ] ↪ 1 cell hidden

- Test the model

[ ] ↪ 4 cells hidden

- Measure the performance using metrics

[ ] ↪ 2 cells hidden

▼ RIDGE

```

8.89624114, 14.76351705, 8.79684319, 12.79443562, 11.75577235,
6.75188629, 11.08634507, 12.21751, 12.27266551, 10.36130386,
12.26162097, 15.75670394, 5.61229575, 5.36044997, 7.44423989,
9.97921082, 11.76463271, 12.11739638, 12.4871439, 13.09112428,
6.78073143, 7.47797529, 6.32351034, 11.14633149, 9.22529064,
11.52155941, 11.03903576, 11.22296093, 10.77208672, 8.7782633,
9.18161649, 9.55926771, 10.52363855, 10.3501496, 8.01031509,
11.44072487, 13.50779466, 7.40142993, 7.55755412, 8.88144185,
4.57893315, 9.39709337, 9.03111482, 11.48432858, 6.56458861,
11.3825303, 9.31036182, 10.34900556, 5.9780508, 5.99561561,
13.16434388, 7.85265927, 11.37483624, 12.43031006, 11.32049169,
6.02664947, 8.81847124, 11.00656069, 11.75424398, 9.48337478,
13.4867177, 8.40227002, 7.64766384, 8.07276202, 11.14796801,
7.6618817, 11.60947783, 4.01231736, 8.69720556, 11.2113068,
10.81846734, 9.72856703, 6.23782707, 8.56416621, 9.10602092,
16.30089766, 12.61879688, 9.81101747, 10.60896902, 11.86558917,
7.97155832, 11.02158753, 11.02573458, 10.08517598, 10.17328851,
6.32373852, 8.78270762, 9.64519434, 9.61061123, 6.35367976,
0.00000000, 12.10147448, 12.10155567, 7.0000701, 11.0027656

```

```
rg.coef_
```

```
array([-0.22437736, -0.70969267, 0.4418621, 0.88051191, 0.66291868,
       -1.70529939, 0.53128287, 1.72113842])
```

```
metrics.r2_score(y_test, rg_pred)
```

```
0.45198921381183166
```

```
np.sqrt(mean_squared_error(y_test, rg_pred))
```

```
2.5477847772203694
```

✓ 0s completed at 10:08 PM

×